

Scilab is not naive

Michael Baudin

January 2010

Abstract

Most of the time, the mathematical formula is directly used in the Scilab source code. But, in many algorithms, some additional work is performed, which takes into account the fact that the computer does not process mathematical real values, but performs computations with their floating point representation. The goal of this article is to show that, in many situations, Scilab is not naive and use algorithms which have been specifically tailored for floating point computers. We analyze in this article the particular case of the quadratic equation, the complex division and the numerical derivatives. In each example, we show that the naive algorithm is not sufficiently accurate, while Scilab implementation is much more robust.

Contents

1	Introduction	2
2	Quadratic equation	4
2.1	Theory	4
2.2	Experiments	5
2.2.1	Massive cancellation	6
2.2.2	Overflow	7
2.3	Explanations	8
2.3.1	Properties of the roots	9
2.3.2	Floating-Point implementation : overview	9
2.3.3	Floating-Point implementation : fixing massive cancellation	10
2.3.4	Floating-Point implementation : fixing overflow problems	11
2.4	References	13
3	Numerical derivatives	13
3.1	Theory	14
3.2	Experiments	14
3.3	Explanations	16
3.3.1	Floating point implementation	16
3.3.2	Robust algorithm	17
3.4	One more step	18
3.5	References	19

4	Complex division	20
4.1	Theory	20
4.2	Experiments	20
4.3	Explanations	22
4.3.1	Algebraic computations	22
4.3.2	Smith's method	24
4.4	One more step	25
4.5	References	28
5	Conclusion	29
6	Acknowledgments	30
7	Appendix	30
7.1	Why 0.1 is rounded	30
7.2	Why $\sin(\pi)$ is rounded	32
7.3	One more step	33
	Bibliography	34
	Index	36

1 Introduction

As a practical example of the problem considered in this document, consider the following numerical experiments. The following session is an example of a Scilab session, where we compute the real number 0.1 by two different, but mathematically equivalent ways.

```
-->format(25)
-->0.1
ans =
    0.10000000000000000055511
-->1.0-0.9
ans =
    0.09999999999999999777955
-->0.1 == 1.0 - 0.9
ans =
F
```

I guess that for a person who has never heard of these problems, this experiment may be a shock. To get things clearer, let's check that the sinus function is also approximated in the sense that the value of $\sin(\pi)$ is *not exactly* zero.

```
-->format(25)
-->sin(0.0)
ans =
    0.
-->sin(%pi)
ans =
    0.00000000000000001224647
```

With symbolic computation systems, such as Maple[28], Mathematica[37] or Maxima[2] for example, the calculations are performed with abstract mathematical symbols. Therefore, there is no loss of accuracy, as long as no numerical evaluation is performed. If a numerical solution is required as a rational number of the form p/q where p and q are integers and $q \neq 0$, there is still no loss of accuracy. On the other hand, in numerical computing systems, such as Scilab[7], Matlab[29] or Octave[3] for example, the computations are performed with floating point numbers. When a numerical value is stored, it is generally associated with a rounding error.

The difficulty of numerical computations is generated by the fact that, while the mathematics treat with *real* numbers, the computer deals with their *floating point representations*. This is the difference between the *naive*, mathematical, approach, and the *numerical*, floating-point, implementation.

In this article, we will not present the floating point arithmetic in detail. Instead, we will show examples of floating point issues by using the following algebraic and experimental approach.

1. First, we will derive the basic theory of a mathematical formula.
2. Then, we will implement it in Scilab and compare with the result given by the equivalent function provided by Scilab. As we will see, some particular cases do not work well with our formula, while the Scilab function computes a correct result.
3. Finally, we will analyze the *reasons* of the differences.

Our numerical experiments will be based on Scilab.

In order to measure the accuracy of the results, we will use two different measures of error: the relative error and the absolute error[19]. Assume that $x_c \in \mathbb{R}$ is a computed value and $x_e \in \mathbb{R}$ is the expected (exact) value. We are looking for a measure of the *difference* between these two real numbers. Most of the time, we use the relative error

$$e_r = \frac{|x_c - x_e|}{|x_e|}, \quad (1)$$

where we assume that $x_e \neq 0$. The relative error e_r is linked with the number of significant digits in the computed value x_c . For example, if the relative error $e_r = 10^{-6}$, then the number of significant digits is 6.

When the expected value is zero, the relative error cannot be computed, and we then use instead the absolute error

$$e_a = |x_c - x_e|. \quad (2)$$

A practical way of checking the expected result of a computation is to compare the formula computed "by hand" with the result produced by a symbolic tool. Recently, Wolfram has launched the <http://www.wolframalpha.com> website which let us access to Mathematica with a classical web browser. Many examples in this document have been validated with this tool.

In the following, we make a brief overview of floating point numbers used in Scilab. Real variables in Scilab are stored in *double precision* floating point variables.

Indeed, Scilab uses the IEEE 754 standard so that real variables are stored with 64 bits floating point numbers, called *doubles*. The floating point number associated with a given $x \in \mathbb{R}$ will be denoted by $fl(x)$.

While the real numbers form a continuum, floating point numbers are both finite and bounded. Not all real numbers can be represented by a floating point number. Indeed, there is an infinite number of reals, while there is a finite number of floating point numbers. In fact, there are, at most, 2^{64} different 64 bits floating point numbers. This leads to *roundoff*, *underflow* and *overflow*.

The double floating point numbers are associated with a machine epsilon equal to 2^{-52} , which is approximately equal to 10^{-16} . This parameter is stored in the `%eps` Scilab variable. Therefore, we can expect, at best, approximately 16 significant decimal digits. This parameter does not depend on the machine we use. Indeed, be it a Linux or a Windows system, Scilab uses IEEE doubles. Therefore, the value of the `%eps` variable is always the same in Scilab.

Negative normalized floating point numbers are in the range $[-10^{308}, -10^{-307}]$ and positive normalized floating point numbers are in the range $[10^{-307}, 10^{308}]$. The limits given in the previous intervals are only decimal approximations. Any real number greater than 10^{309} or smaller than -10^{309} is not representable as a double and is stored with the "infinite" value: in this case, we say that an overflow occurred. A real which magnitude is smaller than 10^{-324} is not representable as a double and is stored as a zero: in this case, we say that an underflow occurred.

The outline of this paper is the following. In the first section, we compute the roots of a quadratic equation. In the second section, we compute the numerical derivatives of a function. In the final section, we perform a numerically difficult division, with complex numbers. The examples presented in this introduction are presented in the appendix of this document.

2 Quadratic equation

In this section, we analyze the computation of the roots of a quadratic polynomial. As we shall see, there is a whole *world* from the mathematical formulas to the implementation of such computations. In the first part, we briefly report the formulas which allow to compute the real roots of a quadratic equation with real coefficients. We then present the naive algorithm based on these mathematical formulas. In the second part, we make some experiments in Scilab and compare our naive algorithm with the `roots` Scilab function. In the third part, we analyze why and how floating point numbers must be taken into account when the roots of a quadratic are required.

2.1 Theory

In this section, we present the mathematical formulas which allow to compute the real roots of a quadratic polynomial with real coefficients. We chose to begin by the example of a quadratic equation, because most of us exactly know (or *think* we know) how to solve such an equation with a computer.

Assume that $a, b, c \in \mathbb{R}$ are given coefficients and $a \neq 0$. Consider the following quadratic [4, 1, 5] equation:

$$ax^2 + bx + c = 0, \quad (3)$$

where $x \in \mathbb{R}$ is the unknown.

Let us define by $\Delta = b^2 - 4ac$ the discriminant of the quadratic equation. We consider the mathematical solution of the quadratic equation, depending on the sign of the discriminant $\Delta = b^2 - 4ac$.

- If $\Delta > 0$, there are two real roots:

$$x_- = \frac{-b - \sqrt{\Delta}}{2a}, \quad (4)$$

$$x_+ = \frac{-b + \sqrt{\Delta}}{2a}. \quad (5)$$

- If $\Delta = 0$, there is one double root:

$$x_{\pm} = -\frac{b}{2a}. \quad (6)$$

- If $\Delta < 0$, there are two complex roots:

$$x_{\pm} = \frac{-b}{2a} \pm i \frac{\sqrt{-\Delta}}{2a}. \quad (7)$$

We now consider a simplified algorithm where we only compute the real roots of the quadratic, assuming that $\Delta > 0$. This naive algorithm is presented in figure 1.

input : a, b, c

output: x_-, x_+

$\Delta := b^2 - 4ac$;

$s := \sqrt{\Delta}$;

$x_- := (-b - s)/(2a)$;

$x_+ := (-b + s)/(2a)$;

Algorithm 1: Naive algorithm to compute the real roots of a quadratic equation.

- We assume that $\Delta > 0$.

2.2 Experiments

In this section, we compare our naive algorithm with the `roots` function. We begin by defining a function which naively implements the mathematical formulas. Then we use our naive function on two particular examples. In the first example, we focus on massive cancellation and in the second example, we focus on overflow problems.

The following Scilab function `myroots` is a straightforward implementation of the previous formulas. It takes as input the coefficients of the quadratic, stored in the vector variable `p`, and returns the two roots in the vector `r`.

```

function r=myroots(p)
    c=coeff(p,0);
    b=coeff(p,1);
    a=coeff(p,2);
    r(1)=(-b+sqrt(b^2-4*a*c))/(2*a);
    r(2)=(-b-sqrt(b^2-4*a*c))/(2*a);
endfunction

```

2.2.1 Massive cancellation

We analyze the rounding errors which are appearing when the discriminant of the quadratic equation is such that $b^2 \gg 4ac$. We consider the following quadratic equation

$$\epsilon x^2 + (1/\epsilon)x - \epsilon = 0 \quad (8)$$

with $\epsilon > 0$. For example, consider the special case $\epsilon = 0.0001 = 10^{-4}$. The discriminant of this equation is $\Delta = 1/\epsilon^2 + 4\epsilon^2$.

The two real solutions of the quadratic equation are

$$x_- = \frac{-1/\epsilon - \sqrt{1/\epsilon^2 + 4\epsilon^2}}{2\epsilon}, \quad x_+ = \frac{-1/\epsilon + \sqrt{1/\epsilon^2 + 4\epsilon^2}}{2\epsilon}. \quad (9)$$

These roots are approximated by

$$x_- \approx -1/\epsilon^2, \quad x_+ \approx \epsilon^2, \quad (10)$$

when ϵ is close to zero. We now consider the limit of the two roots when $\epsilon \rightarrow 0$. We have

$$\lim_{\epsilon \rightarrow 0} x_- = -\infty, \quad \lim_{\epsilon \rightarrow 0} x_+ = 0. \quad (11)$$

In the following Scilab script, we compare the roots computed by the `roots` function and the roots computed by our naive function. We begin by creating a polynomial with the `poly` function, which is given the coefficients of the polynomial. Only the positive root $x_+ \approx \epsilon^2$ is considered in this test. Indeed, the x_- root is so that $x_- \rightarrow -\infty$ in both implementations.

```

p=poly([-0.0001 10000.0 0.0001],"x","coeff");
e1 = 1e-8;
roots1 = myroots(p);
r1 = roots1(1);
roots2 = roots(p);
r2 = roots2(1);
error1 = abs(r1-e1)/e1;
error2 = abs(r2-e1)/e1;
printf("Expected : %e\n", e1);
printf("Naive method : %e (error=%e)\n", r1,error1);
printf("Scilab method : %e (error=%e)\n", r2, error2);

```

The previous script produces the following output.

```

Expected : 1.000000e-008
Naive method : 9.094947e-009 (error=9.050530e-002)
Scilab method : 1.000000e-008 (error=1.654361e-016)

```

We see that the naive method produces a root which has no significant digit and a relative error which is 14 orders of magnitude greater than the relative error of the Scilab root.

This behavior is explained by the fact that the expression for the positive root x_+ given by the equality 5 is numerically evaluated as following. We first consider how the discriminant $\Delta = 1/\epsilon^2 + 4\epsilon^2$ is computed. The term $1/\epsilon^2$ is equal to 100000000 and the term $4\epsilon^2$ is equal to 0.00000004. Therefore, the sum of these two terms is equal to 100000000.000000045. Hence, the square root of the discriminant is

$$\sqrt{1/\epsilon^2 + 4\epsilon^2} = 10000.0000000000001818989. \quad (12)$$

As we see, the first digits are correct, but the last digits are subject to rounding errors. When the expression $-1/\epsilon + \sqrt{1/\epsilon^2 + 4\epsilon^2}$ is evaluated, the following computations are performed :

$$-1/\epsilon + \sqrt{1/\epsilon^2 + 4\epsilon^2} = -10000.0 + 10000.0000000000001818989 \quad (13)$$

$$= 0.00000000000018189894035 \quad (14)$$

We see that the result is mainly driven by the cancellation of significant digits.

We may think that the result is extreme, but it is not. For example, consider the case where we reduce further the value of ϵ down to $\epsilon = 10^{-11}$, we get the following output :

```

Expected : 1.000000e-022
Naive method : 0.000000e+000 (error=1.000000e+000)
Scilab method : 1.000000e-022 (error=1.175494e-016)

```

The relative error is this time 16 orders of magnitude greater than the relative error of the Scilab root. There is no significant decimal digit in the result. In fact, the naive implementation computes a false root x_+ even for a value of epsilon equal to $\epsilon = 10^{-3}$, where the relative error is 7 orders of magnitude greater than the relative error produced by the **roots** function.

2.2.2 Overflow

In this section, we analyse the overflow which appears when the discriminant of the quadratic equation is such that $b^2 - 4ac$ is not representable as a double. We consider the following quadratic equation

$$x^2 + (1/\epsilon)x + 1 = 0 \quad (15)$$

with $\epsilon > 0$. We especially consider the case $\epsilon \rightarrow 0$. The discriminant of this equation is $\Delta = 1/\epsilon^2 - 4$. Assume that the discriminant is positive. Therefore, the roots of the quadratic equation are

$$x_- = \frac{-1/\epsilon - \sqrt{1/\epsilon^2 - 4}}{2}, \quad x_+ = \frac{-1/\epsilon + \sqrt{1/\epsilon^2 - 4}}{2}. \quad (16)$$

These roots are approximated by

$$x_- \approx -1/\epsilon, \quad x_+ \approx -\epsilon, \quad (17)$$

when ϵ is close to zero. We now consider the limit of the two roots when $\epsilon \rightarrow 0$. We have

$$\lim_{\epsilon \rightarrow 0} x_- = -\infty, \quad \lim_{\epsilon \rightarrow 0} x_+ = 0^-. \quad (18)$$

To create a difficult case, we search ϵ so that $1/\epsilon^2 > 10^{308}$, because we know that 10^{308} is the maximum representable double precision floating point number. Therefore, we expect that something should go wrong in the computation of the expression $\sqrt{1/\epsilon^2 - 4}$. We choose $\epsilon = 10^{-155}$.

In the following script, we compare the roots computed by the `roots` function and our naive implementation.

```
e=1.e-155
a = 1;
b = 1/e;
c = 1;
p=poly([c b a],"x","coeff");
expected = [-e;-1/e];
roots1 = myroots(p);
roots2 = roots(p);
error1 = abs(roots1-expected)/norm(expected);
error2 = abs(roots2-expected)/norm(expected);
printf("Expected : %e %e\n", expected(1),expected(2));
printf("Naive method : %e %e (error=%e %e)\n", ...
    roots1(1),roots1(2), error1(1),error1(2));
printf("Scilab method : %e %e (error=%e %e)\n", ...
    roots2(1),roots2(2), error2(1),error2(2));
```

The previous script produces the following output.

```
Expected : -1.000000e-155 -1.000000e+155
Naive method : Inf Inf (error=Nan Nan)
Scilab method : -1.000000e-155 -1.000000e+155
               (error=0.000000e+000 0.000000e+000)
```

In this case, the discriminant $\Delta = b^2 - 4ac$ has been evaluated as $1/\epsilon^2 - 4$, which is approximately equal to 10^{310} . This number cannot be represented in a double precision floating point number. It therefore produces the IEEE Infinite number, which is displayed by Scilab as `Inf`. The Infinite number is associated with an algebra and functions can perfectly take this number as input. Therefore, when the square root function must compute $\sqrt{\Delta}$, it produces again `Inf`. This number is then propagated into the final roots.

2.3 Explanations

In this section, we suggest robust methods to compute the roots of a quadratic equation.

The methods presented in this section are extracted from the *quad* routine of the *RPOLY* algorithm by Jenkins and Traub [22, 21]. This algorithm is used by Scilab in the `roots` function, where a special case is used when the degree of the equation is equal to 2, i.e. a quadratic equation.

2.3.1 Properties of the roots

In this section, we present elementary results, which will be used in the derivation of robust floating point formulas of the roots of the quadratic equation.

Let us assume that the quadratic equation 3, with real coefficients $a, b, c \in \mathbb{R}$ and $a > 0$ has a positive discriminant $\Delta = b^2 - 4ac$. Therefore, the two real roots of the quadratic equation are given by the equations 4 and 5. We can prove that the sum and the product of the roots satisfy the equations

$$x_- + x_+ = \frac{-b}{a}, \quad x_- x_+ = \frac{c}{a}. \quad (19)$$

Therefore, the roots are the solution of the normalized quadratic equation

$$x^2 - (x_- + x_+)x + x_- x_+ = 0. \quad (20)$$

Another transformation leads to an alternative form of the roots. Indeed, the original quadratic equation can be written as a quadratic equation of the unknown $1/x$. Consider the quadratic equation 3 and divide it by $1/x^2$, assuming that $x \neq 0$. This leads to the equation

$$c(1/x)^2 + b(1/x) + a = 0, \quad (21)$$

where we assume that $x \neq 0$. The two real roots of the quadratic equation 21 are

$$x_- = \frac{2c}{-b + \sqrt{b^2 - 4ac}}, \quad (22)$$

$$x_+ = \frac{2c}{-b - \sqrt{b^2 - 4ac}}. \quad (23)$$

The expressions 22 and 23 can also be derived directly from the equations 4 and 5. For that purpose, it suffices to multiply their numerator and denominator by $-b + \sqrt{b^2 - 4ac}$.

2.3.2 Floating-Point implementation : overview

The numerical experiments presented in sections 2.2.1 and 2.2.2 suggest that the floating point implementation must deal with two different problems:

- massive cancellation when $b^2 \gg 4ac$ because of the cancellation of the terms $-b$ and $\pm\sqrt{b^2 - 4ac}$ which may have opposite signs,
- overflow in the computation of the square root of the discriminant $\sqrt{\pm(b^2 - 4ac)}$ when $b^2 - 4ac$ is not representable as a floating point number.

The cancellation problem occurs only when the discriminant is positive, i.e. only when there are two real roots. Indeed, the cancellation will not appear when $\Delta < 0$, since the complex roots do not use the sum $-b \pm \sqrt{b^2 - 4ac}$. When $\Delta = 0$, the double real root does not cause any trouble. Therefore, we must take into account for the cancellation problem only in the equations 4 and 5.

On the other hand, the overflow problem occurs whatever the sign of the discriminant but does not occur when $\Delta = 0$. Therefore, we must take into account for this problem in the equations 4, 5 and 7. In section 2.3.3, we focus on the cancellation error while the overflow problem is addressed in section 2.3.4.

2.3.3 Floating-Point implementation : fixing massive cancellation

In this section, we present the computation of the roots of a quadratic equation with protection against massive cancellation.

When the discriminant Δ is positive, the massive cancellation problem can be split in two cases:

- if $b < 0$, then $-b - \sqrt{b^2 - 4ac}$ may suffer of massive cancellation because $-b$ is positive and $-\sqrt{b^2 - 4ac}$ is negative,
- if $b > 0$, then $-b + \sqrt{b^2 - 4ac}$ may suffer of massive cancellation because $-b$ is negative and $\sqrt{b^2 - 4ac}$ is positive.

Therefore,

- if $b > 0$, we should use the expression $-b - \sqrt{b^2 - 4ac}$,
- if $b < 0$, we should use the expression $-b + \sqrt{b^2 - 4ac}$.

The solution consists in a combination of the following expressions of the roots given by, on one hand the equations 4 and 5, and, on the other hand the equations 22 and 23. We pick the formula so that the sign of b is the same as the sign of the square root. The following choice allow to solve the massive cancellation problem:

- if $b < 0$, then compute x_- from 22, else (if $b > 0$), compute x_- from 4,
- if $b < 0$, then compute x_+ from 5, else (if $b > 0$), compute x_+ from 23.

We can also consider the modified Fagnano formulas

$$x_1 = -\frac{2c}{b + \operatorname{sgn}(b)\sqrt{b^2 - 4ac}}, \quad (24)$$

$$x_2 = -\frac{b + \operatorname{sgn}(b)\sqrt{b^2 - 4ac}}{2a}, \quad (25)$$

where the sign function is defined by

$$\operatorname{sgn}(b) = \begin{cases} 1, & \text{if } b \geq 0, \\ -1, & \text{if } b < 0. \end{cases} \quad (26)$$

The roots $x_{1,2}$ correspond to the roots $x_{+,-}$. Indeed, on one hand, if $b < 0$, $x_1 = x_-$ and if $b > 0$, $x_1 = x_+$. On the other hand, if $b < 0$, $x_2 = x_+$ and if $b > 0$, $x_2 = x_-$.

Moreover, we notice that the division by two (and the multiplication by 2) is exact with floating point numbers so these operations cannot be a source of problem. But it is interesting to use $b/2$, which involves only one division, instead of the three multiplications $2 * c$, $2 * a$ and $4 * a * c$. This leads to the following expressions of the real roots

$$x_1 = -\frac{c}{(b/2) + \operatorname{sgn}(b)\sqrt{(b/2)^2 - ac}}, \quad (27)$$

$$x_2 = -\frac{(b/2) + \operatorname{sgn}(b)\sqrt{(b/2)^2 - ac}}{a}. \quad (28)$$

Therefore, the two real roots can be computed by the following sequence of computations:

$$b' := b/2, \quad \Delta' := b'^2 - ac, \quad (29)$$

$$h := -\left(b' + \text{sgn}(b)\sqrt{\Delta'}\right) \quad (30)$$

$$x_1 := \frac{c}{h}, \quad x_2 := \frac{h}{a}. \quad (31)$$

In the case where the discriminant $\Delta' := b'^2 - ac$ is negative, the two complex roots are

$$x_1 = -\frac{b'}{a} - i\frac{\sqrt{ac - b'^2}}{a}, \quad x_2 = -\frac{b'}{a} + i\frac{\sqrt{ac - b'^2}}{a}. \quad (32)$$

A more robust algorithm, based on the previous analysis is presented in figure 2. By comparing 1 and 2, we can see that the algorithms are different in many points.

2.3.4 Floating-Point implementation : fixing overflow problems

The remaining problem is to compute the square root of the discriminant $\sqrt{\pm(b'^2 - ac)}$ without creating unnecessary overflows. In order to simplify the discussion, we focus on the computation of $\sqrt{b'^2 - ac}$.

Obviously, the problem occur for large values of b' . Notice that a (very) small improvement has already been done. Indeed, we have the inequality $|b'| = |b|/2 < |b|$ so that overflows are twice less likely to occur. The current upper bound for $|b'|$ is 10^{154} , which is associated with $b'^2 \leq 10^{308}$, the maximum double value before overflow. The goal is therefore to increase the possible range of values of b' without generating unnecessary overflows.

Consider the situation when b' is large in magnitude with respect to a and c . In that case, notice that we first square b' to get b'^2 and then compute the square root $\sqrt{b'^2 - ac}$. Hence, we can factor the expression by b'^2 and move this term outside the square root, which makes the term $|b'|$ appear. This method allows to compute the expression $\sqrt{b'^2 - ac}$, without squaring b' when it is not necessary.

In the general case, we use the fact that the term $b'^2 - ac$ can be evaluated with the two following equivalent formulas:

$$b'^2 - ac = b'^2 [1 - (a/b')(c/b')], \quad (33)$$

$$b'^2 - ac = c [b'(b'/c) - a]. \quad (34)$$

The goal is then to compute the square root $s = \sqrt{b'^2 - ac}$.

- If $|b'| > |c| > 0$, then the equation 33 involves the expression $1 - (a/b')(c/b')$. The term $1 - (a/b')(c/b')$ is so that no overflow is possible since $|c/b'| < 1$ (the overflow problem occurs only when b is large in magnitude with respect to both a and c). In this case, we use the expression

$$e = 1 - (a/b')(c/b'), \quad (35)$$

```

input :  $a, b, c$ 
output:  $x_-^R, x_-^I, x_+^R, x_+^I$ 
if  $a = 0$  then
  if  $b = 0$  then
     $x_-^R := 0, x_-^I := 0;$ 
     $x_+^R := 0, x_+^I := 0;$ 
  else
     $x_-^R := -c/b, x_-^I := 0;$ 
     $x_+^R := 0, x_+^I := 0;$ 
  end
else
   $b' := b/2;$ 
   $\Delta := b'^2 - ac;$ 
  if  $\Delta < 0$  then
     $s := \sqrt{-\Delta};$ 
     $x_-^R := -b'/a, x_-^I := -s/a;$ 
     $x_+^R := x_-^R, x_+^I := -x_-^I;$ 
  else if  $\Delta = 0$  then
     $x_- := -b'/a, x_-^I := 0;$ 
     $x_+ := x_-, x_+^I := 0;$ 
  else
     $s := \sqrt{\Delta};$ 
    if  $b > 0$  then
       $g := 1;$ 
    else
       $g := -1;$ 
    end
     $h := -(b' + g * s);$ 
     $x_-^R := c/h, x_-^I := 0;$ 
     $x_+^R := h/a, x_+^I := 0;$ 
  end
end

```

Algorithm 2: A more robust algorithm to compute the roots of a quadratic equation. This algorithm takes as input arguments the real coefficients a, b, c and returns the real and imaginary parts of the two roots, i.e. returns $x_-^R, x_-^I, x_+^R, x_+^I$.

and compute

$$s = \pm |b'| \sqrt{|e|}. \quad (36)$$

In the previous equation, we use the sign $+$ when e is positive and the sign $-$ when e is negative.

- If $|c| > |b'| > 0$, then the equation 34 involves the expression $b'(b'/c) - a$. The term $b'(b'/c) - a$ should limit the possible overflows since $|b'/c| < 1$. This implies that $|b'(b'/c)| < |b'|$. (There is still a possibility of overflow, for example in the case where $b'(b'/c)$ is near, but under, the overflow limit and a is large.) Therefore, we use the expression

$$e = b'(b'/c) - a, \quad (37)$$

and compute

$$s = \pm \sqrt{|c|} \sqrt{|e|}. \quad (38)$$

In the previous equation, we use the sign $+$ when e is positive and the sign $-$ when e is negative.

In both equations 36 and 38, the parenthesis must be strictly used. This property is ensured by the IEEE standard and by the Scilab language. This normalization method are similar to the one used by Smith in the algorithm for the division of complex numbers [40] and which will be reviewed in the next section.

2.4 References

The 1966 technical report by G. Forsythe [12] presents the floating point system and the possible large error in using mathematical algorithms blindly. An accurate way of solving a quadratic is outlined. A few general remarks are made about computational mathematics.

The 1991 paper by Goldberg [18] is a general presentation of the floating point system and its consequences. It begins with background on floating point representation and rounding errors, continues with a discussion of the IEEE floating point standard and concludes with examples of how computer system builders can better support floating point. The section 1.4, "Cancellation" specifically consider the computation of the roots of a quadratic equation.

We can also read the numerical experiments performed by Nievergelt in [34].

The Numerical Recipes [35], chapter 5, section 5.6, "Quadratic and Cubic Equations" present the elementary theory for a floating point implementation of the quadratic and cubic equations.

Other references include William Kahan [24].

3 Numerical derivatives

In this section, we analyze the computation of the numerical derivative of a given function.

In the first part, we briefly report the first order forward formula, which is based on the Taylor theorem. We then present the naive algorithm based on these mathematical formulas. In the second part, we make some experiments in Scilab and compare our naive algorithm with the `derivative` Scilab function. In the third part, we analyze why and how floating point numbers must be taken into account when we compute numerical derivatives.

3.1 Theory

The basic result is the Taylor formula with one variable [20]. Assume that $x \in \mathbb{R}$ is a given number and $h \in \mathbb{R}$ is a given step. Assume that $f : \mathbb{R} \rightarrow \mathbb{R}$ is a two times continuously differentiable function. Therefore,

$$f(x+h) = f(x) + hf'(x) + \frac{h^2}{2}f''(x) + \mathcal{O}(h^3). \quad (39)$$

We immediately get the forward difference which approximates the first derivate at order 1

$$f'(x) = \frac{f(x+h) - f(x)}{h} + \frac{h}{2}f''(x) + \mathcal{O}(h^2). \quad (40)$$

The naive algorithm to compute the numerical derivate of a function of one variable is presented in figure 3.

input : x, h
output: $f'(x)$
 $f'(x) := (f(x+h) - f(x))/h$;
Algorithm 3: Naive algorithm to compute the numerical derivative of a function of one variable.

3.2 Experiments

The following Scilab function `myfprime` is a straightforward implementation of the previous algorithm.

```
function fp = myfprime(f,x,h)
    fp = (f(x+h) - f(x))/h;
endfunction
```

In our experiments, we will compute the derivatives of the square function $f(x) = x^2$, which is $f'(x) = 2x$. The following Scilab function `myfunction` computes the square function.

```
function y = myfunction (x)
    y = x*x;
endfunction
```

The (naive) idea is that the computed relative error is small when the step h is small. Because *small* is not a priori clear, we take $h = 10^{-16}$ as a "good" candidate for a *small* double.

The `derivative` function allows to compute the Jacobian and the Hessian matrix of a given function. Moreover, we can use formulas of order 1, 2 or 4. The `derivative` function has been designed by Rainer von Seggern and Bruno Pinçon. The order 1 formula is the forward numerical derivative that we have already presented.

In the following script, we compare the computed relative error produced by our naive method with step $h = 10^{-16}$ and the `derivative` function with default optimal step. We compare the two methods for the point $x = 1$.

```
x = 1.0;
fpref = derivative(myfunction,x,order=1);
e = abs(fpref-2.0)/2.0;
mprintf("Scilab f'='%e, error=%e\n", fpref,e);
h = 1.e-16;
fp = myfprime(myfunction,x,h);
e = abs(fp-2.0)/2.0;
mprintf("Naive f'='%e, h=%e, error=%e\n", fp,h,e);
```

The previous script produces the following output.

```
Scilab f'=2.000000e+000, error=7.450581e-009
Naive f'=0.000000e+000, h=1.000000e-016, error=1.000000e+000
```

Our naive method seems to be inaccurate and has no significant decimal digit. The Scilab function, instead, has 9 significant digits.

Since our faith is based on the truth of the mathematical theory, some deeper experiments must be performed. We make the following numerical experiment: we take the initial step $h = 1.0$ and divide h by 10 at each step of a loop made of 20 iterations.

```
x = 1.0;
fpref = derivative(myfunction,x,order=1);
e = abs(fpref-2.0)/2.0;
mprintf("Scilab f'='%e, error=%e\n", fpref,e);
h = 1.0;
for i=1:20
    h=h/10.0;
    fp = myfprime(myfunction,x,h);
    e = abs(fp-2.0)/2.0;
    mprintf("Naive f'='%e, h=%e, error=%e\n", fp,h,e);
end
```

The previous script produces the following output.

```
Scilab f'=2.000000e+000, error=7.450581e-009
Naive f'=2.100000e+000, h=1.000000e-001, error=5.000000e-002
Naive f'=2.010000e+000, h=1.000000e-002, error=5.000000e-003
Naive f'=2.001000e+000, h=1.000000e-003, error=5.000000e-004
Naive f'=2.000100e+000, h=1.000000e-004, error=5.000000e-005
Naive f'=2.000010e+000, h=1.000000e-005, error=5.000007e-006
Naive f'=2.000001e+000, h=1.000000e-006, error=4.999622e-007
Naive f'=2.000000e+000, h=1.000000e-007, error=5.054390e-008
Naive f'=2.000000e+000, h=1.000000e-008, error=6.077471e-009
Naive f'=2.000000e+000, h=1.000000e-009, error=8.274037e-008
Naive f'=2.000000e+000, h=1.000000e-010, error=8.274037e-008
Naive f'=2.000000e+000, h=1.000000e-011, error=8.274037e-008
Naive f'=2.000178e+000, h=1.000000e-012, error=8.890058e-005
Naive f'=1.998401e+000, h=1.000000e-013, error=7.992778e-004
```

```

Naive f'=1.998401e+000, h=1.000000e-014, error=7.992778e-004
Naive f'=2.220446e+000, h=1.000000e-015, error=1.102230e-001
Naive f'=0.000000e+000, h=1.000000e-016, error=1.000000e+000
Naive f'=0.000000e+000, h=1.000000e-017, error=1.000000e+000
Naive f'=0.000000e+000, h=1.000000e-018, error=1.000000e+000
Naive f'=0.000000e+000, h=1.000000e-019, error=1.000000e+000
Naive f'=0.000000e+000, h=1.000000e-020, error=1.000000e+000

```

We see that the relative error begins by decreasing, gets to a minimum and then increases. Obviously, the optimum step is approximately $h = 10^{-8}$, where the relative error is approximately $e_r = 6.10^{-9}$. We should not be surprised to see that Scilab has computed a derivative which is near the optimum.

3.3 Explanations

In this section, we make reasonable assumptions for the expression of the total error and compute the optimal step of a forward difference formula. We extend our work to the centered two points formula.

3.3.1 Floating point implementation

The first source of error is obviously the truncation error $E_t(h) = h|f''(x)|/2$, due to the limited Taylor expansion.

The other source of error is generated by the roundoff errors in the function evaluation of the formula $(f(x+h) - f(x))/h$. Indeed, the floating point representation of the function value at point x is

$$fl(f(x)) = (1 + e(x))f(x), \quad (41)$$

where the relative error e depends on the the current point x . We assume here that the relative error e is bounded by the product of a constant $c > 0$ and the machine precision r . Furthermore, we assume here that the constant c is equal to one. We may consider other rounding errors sources, such as the error in the sum $x + h$, the difference $f(x+h) - f(x)$ or the division $(f(x+h) - f(x))/h$. But all these rounding errors can be neglected for they are not, in general, as large as the roundoff error generated by the function evaluation. Hence, the roundoff error associated with the function evaluation is $E_r(h) = r|f(x)|/h$.

Therefore, the total error associated with the forward finite difference is bounded by

$$E(h) = \frac{r|f(x)|}{h} + \frac{h}{2}|f''(x)|. \quad (42)$$

The error is then the sum of a term which is a decreasing function of h and a term which an increasing function of h . We consider the problem of finding the step h which minimizes the error $E(h)$. The total error $E(h)$ is minimized when its first derivative is zero. The first derivative of the function E is

$$E'(h) = -\frac{r|f(x)|}{h^2} + \frac{1}{2}|f''(x)|. \quad (43)$$

The second derivative of E is

$$E''(h) = 2 \frac{r|f(x)|}{h^3}. \quad (44)$$

If we assume that $f(x) \neq 0$, then the second derivative $E''(h)$ is strictly positive, since $h > 0$ (i.e. we consider only non-zero steps). Hence, there is only one global solution of the minimization problem. This first derivative is zero if and only if

$$-\frac{r|f(x)|}{h^2} + \frac{1}{2}|f''(x)| = 0 \quad (45)$$

Therefore, the optimal step is

$$\bar{h} = \sqrt{\frac{2r|f(x)|}{|f''(x)|}}. \quad (46)$$

Let us make the additional assumption

$$\frac{2|f(x)|}{|f''(x)|} \approx 1. \quad (47)$$

Then the optimal step is

$$\bar{h} = \sqrt{r}, \quad (48)$$

where the error is

$$E(\bar{h}) = 2\sqrt{r}. \quad (49)$$

With double precision floating point numbers, we have $r = 10^{-16}$ and we get $\bar{h} = 10^{-8}$ and $E(\bar{h}) = 2.10^{-8}$. Under our assumptions on f and on the form of the total error, this is the minimum error which is achievable with a forward difference numerical derivate.

We can extend the previous method to the first derivate computed by a centered 2 points formula. We can prove that

$$f'(x) = \frac{f(x+h) - f(x-h)}{2h} + \frac{h^2}{6}f'''(x) + \mathcal{O}(h^3). \quad (50)$$

We can apply the same method as previously and, under reasonable assumptions on f and the form of the total error, we get that the optimal step is $h = r^{1/3}$, which corresponds to the total error $E = 2r^{2/3}$. With double precision floating point numbers, this corresponds to $h \approx 10^{-5}$ and $E \approx 10^{-10}$.

3.3.2 Robust algorithm

A more robust algorithm to compute the numerical derivate of a function of one variable is presented in figure 4.

$$h := \sqrt{r};$$

$$f'(x) := (f(x+h) - f(x))/h;$$

Algorithm 4: A more robust algorithm to compute the numerical derivative of a function of one variable.

3.4 One more step

In this section, we analyze the behavior of the `derivative` function when the point x is either large in magnitude, small or close to zero. We compare these results with the `numdiff` function, which does not use the same step strategy. As we are going to see, both functions performs the same when x is near 1, but performs very differently when x is large or small.

The `derivative` function uses the optimal step based on the theory we have presented. But the optimal step does not solve all the problems that may occur in practice, as we are going to see.

See for example the following Scilab session, where we compute the numerical derivative of $f(x) = x^2$ for $x = 10^{-100}$. The expected result is $f'(x) = 2 \times 10^{-100}$.

```
-->fp = derivative(myfunction,1.e-100,order=1)
fp =
    0.00000000149011611938477
-->fe=2.e-100
fe =
    2.00000000000000000040-100
-->e = abs(fp-fe)/fe
e =
    7.450580596923828243D+91
```

The result does not have any significant digit.

The explanation is that the step is $h = \sqrt{r} \approx 10^{-8}$. Then, the point $x + h$ is computed as $10^{-100} + 10^{-8}$ which is represented by a floating point number which is close to 10^{-8} , because the term 10^{-100} is much smaller than 10^{-8} . Then we evaluate the function, which leads to $f(x+h) = f(10^{-8}) = 10^{-16}$. The result of the computation is therefore $(f(x+h) - f(x))/h = (10^{-16} + 10^{-200})/10^{-8} \approx 10^{-8}$.

That experiment shows that the `derivative` function uses a poor default step h when x is very small.

To improve the accuracy of the computation, we can take the control of the step h . A reasonable solution is to use $h = \sqrt{r}|x|$ so that the step is scaled depending on x . The following script illustrates than method, which produces results with 8 significant digits.

```
-->fp = derivative(myfunction,1.e-100,order=1,h=sqrt(%eps)*1.e-100)
fp =
    2.0000000013099139394-100
-->fe=2.e-100
fe =
    2.00000000000000000040-100
-->e = abs(fp-fe)/fe
e =
    0.0000000065495696770794
```

But when x is exactly zero, the step $h = \sqrt{r}|x|$ cannot work, because it would produce the step $h = 0$, which would generate a division by zero exception. In that case, the step $h = \sqrt{r}$ provides a sufficiently good accuracy.

Another function is available in Scilab to compute the numerical derivatives of a given function, that is `numdiff`. The `numdiff` function uses the step

$$h = \sqrt{r}(1 + 10^{-3}|x|). \quad (51)$$

In the following paragraphs, we analyze why this formula has been chosen. As we are going to check experimentally, this step formula performs better than `derivative` when x is large, but performs equally bad when x is small.

As we can see the following session, the behavior is approximately the same when the value of x is 1.

```
-->fp = numdiff(myfunction,1.0)
fp =
    2.0000000189353417390237
-->fe=2.0
fe =
    2.
-->e = abs(fp-fe)/fe
e =
    9.468D-09
```

The accuracy is slightly decreased with respect to the optimal value 7.450581e-009 which was produced by the `derivative` function. But the number of significant digits is approximately the same, i.e. 9 digits.

The goal of the step used by the `numdiff` function is to produce good accuracy when the value of x is large. In this case, the `numdiff` function produces accurate results, while the `derivative` function performs poorly.

In the following session, we compute the numerical derivative of the function $f(x) = x^2$ at the point $x = 10^{10}$. The expected result is $f'(x) = 2 \cdot 10^{10}$.

```
-->numdiff(myfunction,1.e10)
ans =
    2.000D+10
-->derivative(myfunction,1.e10,order=1)
ans =
    0.
```

We see that the `numdiff` function produces an accurate result while the `derivative` function produces a result which has no significant digit.

The behavior of the two functions when x is close to zero is the same, i.e. both functions produce wrong results. Indeed, when we use the `derivative` function, the step $h = \sqrt{r}$ is too large so that the point x is neglected against the step h . On the other hand, when we use the `numdiff` function, the step $h = \sqrt{r}(1 + 10^{-3}|x|)$ is approximated by $h = \sqrt{r}$ so that it produces the same results as the `derivative` function.

3.5 References

A reference for numerical derivatives is [6], chapter 25. "Numerical Interpolation, Differentiation and Integration" (p. 875). The webpage [39] and the book [35] give

results about the rounding errors.

In order to solve this issue generated by the magnitude of x , more complex methods should be used. Moreover, we did not give the solution of other sources of rounding errors. Indeed, the step $h = \sqrt{r}$ was computed based on assumptions on the rounding error of the function evaluations, where we consider that the constant c is equal to one. This assumption is satisfied only in the ideal case. Furthermore, we make the assumption that the factor $\frac{2|f(x)|}{|f''(x)|}$ is close to one. This assumption is far from being achieved in practical situations, where the function value and its second derivative can vary greatly in magnitude.

Several authors attempted to solve the problems associated with numerical derivatives. A non-exhaustive list of references includes [25, 11, 41, 16].

4 Complex division

In this section, we analyze the problem of the complex division in Scilab. We especially detail the difference between the mathematical, straightforward formula and the floating point implementation. In the first part, we briefly report the formulas which allow to compute the real and imaginary parts of the division of two complex numbers. We then present the naive algorithm based on these mathematical formulas. In the second part, we make some experiments in Scilab and compare our naive algorithm with Scilab's division operator. In the third part, we analyze why and how floating point numbers must be taken into account when the implementation of such division is required.

4.1 Theory

Assume that a, b, c and d are four real numbers. Consider the two complex numbers $a + ib$ and $c + id$, where i is the imaginary number which satisfies $i^2 = -1$. Assume that $c^2 + d^2$ is non zero. We are interested in the complex number $e + fi = \frac{a+ib}{c+id}$ where e and f are real numbers. The formula which allows to compute the real and imaginary parts of the division of these two complex numbers is

$$\frac{a + ib}{c + id} = \frac{ac + bd}{c^2 + d^2} + i \frac{bc - ad}{c^2 + d^2}. \quad (52)$$

So that the real and imaginary parts e and f of the complex number are

$$e = \frac{ac + bd}{c^2 + d^2}, \quad (53)$$

$$f = \frac{bc - ad}{c^2 + d^2}. \quad (54)$$

The naive algorithm for the computation of the complex division is presented in figure 5.

4.2 Experiments

The following Scilab function `naive` is a straightforward implementation of the previous formulas. It takes as input the complex numbers a and b , represented by their

input : a, b, c, d
output: e, f
 $den := c^2 + d^2;$
 $e := (ac + bd)/den;$
 $f := (bc - ad)/den;$

Algorithm 5: Naive algorithm to compute the complex division. The algorithm takes as input the real and imaginary parts a, b, c, d of the two complex numbers and returns e and f , the real and imaginary parts of the division.

real and imaginary parts **a**, **b**, **c** and **d**. The function **naive** returns the complex number represented by its real and imaginary parts **e** and **f**.

```

function [e,f] = naive (a , b , c , d )
    den = c * c + d * d;
    e = (a * c + b * d) / den;
    f = (b * c - a * d) / den;
endfunction

```

Consider the complex division

$$\frac{1 + i2}{3 + i4} = \frac{11}{25} + i\frac{2}{25} = 0.44 + i0.08. \quad (55)$$

We check our result with Wolfram Alpha[38], with the input `"(1+i*2)/(3+i*4)"`. In the following script, we check that there is no obvious bug in the naive implementation.

```

--> [e f] = naive ( 1.0 , 2.0 , 3.0 , 4.0 )
f =
    0.08
e =
    0.44
--> (1.0 + %i * 2.0 )/(3.0 + %i * 4.0 )
ans =
    0.44 + 0.08i

```

The results of the **naive** function and the division operator are the same, which makes us confident that our implementation is correct.

Now that we are confident, we make the following numerical experiment involving a large number. Consider the complex division

$$\frac{1 + i}{1 + i10^{307}} \approx 1.0000000000000000 \cdot 10^{-307} - i1.0000000000000000 \cdot 10^{-307}, \quad (56)$$

which is accurate to the displayed digits. We check our result with Wolfram Alpha[38], with the input `"(1 + i)/(1 + i * 10^307)"`. In fact, there are more that 300 zeros following the leading 1, so that the previous approximation is very accurate. The following Scilab session compares the naive implementation and Scilab's division operator.

```

--> [e f] = naive ( 1.0 , 1.0 , 1.0 , 1.e307 )
f =
    0.
e =

```

```

0.
--> (1.0 + %i * 1.0)/(1.0 + %i * 1.e307)
ans =
1.000-307 - 1.000-307i

```

In the previous case, the naive implementation does not produce any correct digit!

The last test involves small numbers in the denominator of the complex fraction. Consider the complex division

$$\frac{1+i}{10^{-307} + i10^{-307}} = \frac{1+i}{10^{-307}(1+i)} = 10^{307}. \quad (57)$$

In the following session, the first statement `ieee(2)` configures the IEEE system so that Inf and Nan numbers are generated instead of Scilab error messages.

```

-->ieee(2);
-->[e f] = naive ( 1.0 , 1.0 , 1.e-307 , 1.e-307 )
f =
Nan
e =
Inf
-->(1.0 + %i * 1.0)/(1.e-307 + %i * 1.e-307)
ans =
1.000+307

```

We see that the naive implementation generates the IEEE numbers Nan and Inf, while the division operator produces the correct result.

4.3 Explanations

In this section, we analyze the reason why the naive implementation of the complex division leads to inaccurate results. In the first section, we perform algebraic computations and shows the problems of the naive formulas. In the second section, we present the Smith's method.

4.3.1 Algebraic computations

In this section, we analyze the results produced by the second and third tests in the previous numerical experiments. We show that the intermediate numbers which appear are not representable as double precision floating point numbers.

Let us analyze the second complex division [56](#). We are going to see that this division is associated with an IEEE overflow. We have $a = 1$, $b = 1$, $c = 1$ and $d = 10^{307}$. By the equations [53](#) and [54](#), we have

$$den = c^2 + d^2 = 1^2 + (10^{307})^2 \quad (58)$$

$$= 1 + 10^{614} \approx 10^{614}, \quad (59)$$

$$e = (ac + bd)/den = (1 * 1 + 1 * 10^{307})/10^{614}, \quad (60)$$

$$\approx 10^{307}/10^{614} \approx 10^{-307}, \quad (61)$$

$$f = (bc - ad)/den = (1 * 1 - 1 * 10^{307})/10^{614} \quad (62)$$

$$\approx -10^{307}/10^{614} \approx -10^{-307}. \quad (63)$$

We see that both the input numbers a, b, c, d are representable and the output numbers $e = 10^{-307}$ and $f = -10^{-307}$ are representable as double precision floating point numbers. We now focus on the floating point representation of the intermediate expressions. We have

$$fl(den) = fl(10^{614}) = Inf, \quad (64)$$

because 10^{614} is not representable as a double precision number. Indeed, the largest positive double is 10^{308} . The IEEE Inf floating point number stands for Infinity and is associated with an overflow. The Inf floating point number is associated with an algebra which defines that $1/Inf = 0$. This is consistent with mathematical limit of the function $1/x$ when $x \rightarrow \infty$. Then, the e and f terms are computed as

$$fl(e) = fl((ac + bd)/den) = fl((1 * 1 + 1 * 10^{307})/Inf) = fl(10^{307}/Inf) = 0 \quad (65)$$

$$fl(f) = fl((bc - ad)/den) = fl((1 * 1 - 1 * 10^{307})/Inf) = fl(-10^{307}/Inf) = 0 \quad (66)$$

Hence, the result is computed without any significant digit, even though both the input and the output numbers are all representable as double precision floating point numbers.

Let us analyze the second complex division 57. We are going to see that this division is associated with an IEEE underflow. We have $a = 1$, $b = 1$, $c = 10^{-307}$ and $d = 10^{-307}$. We now use the equations 53 and 54, which leads to:

$$den = c^2 + d^2 = (10^{-307})^2 + (10^{-307})^2 \quad (67)$$

$$= 10^{-614} + 10^{-614} = 2.10^{-614}, \quad (68)$$

$$e = (ac + bd)/den = (1 * 10^{-307} + 1 * 10^{-307})/(2.10^{-614}) \quad (69)$$

$$= (2.10^{-307})/(2.10^{-614}) = 10^{307}, \quad (70)$$

$$f = (bc - ad)/den = (1 * 10^{-307} - 1 * 10^{-307})/(2.10^{-614}) \quad (71)$$

$$= 0/10^{-614} = 0. \quad (72)$$

With double precision floating point numbers, the computation is not performed this way. The positive terms which are smaller than 10^{-324} are too small to be representable in double precision and are represented by 0 so that an underflow occurs. This leads to

$$fl(den) = fl(c^2 + d^2) = fl(10^{-614} + 10^{-614}) \quad (73)$$

$$= 0, \quad (74)$$

$$fl(e) = fl((ac + bd)/den) = fl((1 * 10^{-307} + 1 * 10^{-307})/(2.10^{-614})) \quad (75)$$

$$= fl(2.10^{-307}/0) = Inf, \quad (76)$$

$$fl(f) = fl((bc - ad)/den) = fl((1 * 10^{-307} - 1 * 10^{-307})/0) \quad (77)$$

$$= fl(0/0) = NaN. \quad (78)$$

The two previous examples shows that, even if both the input and output numbers are representable as floating point numbers, the intermediate expressions may generate numbers which may not be representable as floating point numbers. Hence, a naive implementation can lead to inaccurate results. In the next section, we present a method which allows to cure most problems generated by the complex division.

4.3.2 Smith's method

In this section, we analyze Smith's method, which allows to produce an accurate division of two complex numbers. We present the detailed steps of this modified algorithm in the particular cases that we have presented.

In Scilab, the algorithm which allows to perform the complex division is done by the `wwdiv` routine, which implements Smith's method [40]. This implementation is due to Bruno Pinçon. Smith's algorithm is based on normalization, which allow to perform the complex division even if the input terms are large or small.

The starting point of the method is the mathematical definition 52, which is reproduced here for simplicity

$$\frac{a + ib}{c + id} = e + if = \frac{ac + bd}{c^2 + d^2} + i \frac{bc - ad}{c^2 + d^2}. \quad (79)$$

Smith's method is based on the rewriting of this formula in two different, but mathematically equivalent, formulas. We have seen that the term $c^2 + d^2$ may generate overflows or underflows. This is caused by intermediate expressions which magnitudes are larger than necessary. The previous numerical experiments suggest that, provided that we had simplified the calculation, the intermediate expressions would not have been unnecessary large.

Consider the term $e = \frac{ac+bd}{c^2+d^2}$ in the equation 79 and assume that $c \neq 0$ and $d \neq 0$. Let us assume that c is large in magnitude with respect to d , i.e. $|d| \ll |c|$. This implies $d/c \leq 1$. We see that the denominator $c^2 + d^2$ squares the number c , which also appears in the numerator. Therefore, we multiply both the numerator and the denominator by $1/c$. If we express e as $\frac{a+b(d/c)}{c+d(d/c)}$, which is mathematically equivalent, we see that there is no more squaring of c . Hence, overflows are less likely to occur in the denominator, since $|d(d/c)| = |d||d/c| \leq |d|$. That is, there is no growth in the magnitude of the terms involved in the computation of the product $d(d/c)$. Similarly, overflows are less likely to occur in the numerator, since $|b(d/c)| = |b||d/c| \leq |b|$. In the opposite case where d is large with respect to c , i.e. $d \gg c$, we could divide the numerator and the denominator by $1/d$. This leads to the formulas

$$\frac{a + ib}{c + id} = \frac{a + b(d/c)}{c + d(d/c)} + i \frac{b - a(d/c)}{c + d(d/c)}, \quad (80)$$

$$= \frac{a(c/d) + b}{c(c/d) + d} + i \frac{b(c/d) - a}{c(c/d) + d}. \quad (81)$$

The previous equations can be simplified as

$$\frac{a + ib}{c + id} = \frac{a + br}{c + dr} + i \frac{b - ar}{c + dr}, \quad r = d/c, \text{ if } c \geq d, \quad (82)$$

$$= \frac{ar + b}{cr + d} + i \frac{br - a}{cr + d}, \quad r = c/d, \text{ if } d \geq c. \quad (83)$$

The following `smith` function implements Smith's method in the Scilab language.

```
function [e,f] = smith ( a , b , c , d )
    if ( abs(d) <= abs(c) ) then
```



```

    r = d/c;
    den = c + d * r;
    e = (a + b * r) / den;
    f = (b - a * r) / den;
else
    r = c/d;
    den = c * r + d;
    e = (a * r + b) / den;
    f = (b * r - a) / den;
end
endfunction

```

We now check that Smith's method performs very well for the difficult complex division that we met earlier in this chapter.

Let us analyze the second complex division 56. We have $a = 1$, $b = 1$, $c = 1$ and $d = 10^{307}$. For this division, Smith's method is the following.

```

if ( |1.e307| <= |1| ) > test false
else
    r = c/d = 1 / 1.e307 = 1.e-307
    den = c * r + d = 1 * 1.e-307 + 1.e307 = 1.e307
    e = (a * r + b)/den = (1 * 1.e-307 + 1) / 1.e307 = 1 / 1.e307
      = 1.e-307
    f = (b * r - a)/den = (1 * 1.e-307 - 1) / 1.e307 = -1 / 1.e307
      = -1.e-308

```

We see that, while the naive division generated an overflow, Smith's method produces the correct result.

Let us analyze the second complex division 57. We have $a = 1$, $b = 1$, $c = 10^{-307}$ and $d = 10^{-307}$.

```

if ( |1.e-307| <= |1.e-307| ) > test true
    r = d/c = 1.e-307 / 1.e-307 = 1
    den = c + d * r = 1.e-307 + 1e-307 * 1 = 2.e-307
    e = (a + b * r) / den = (1 + 1 * 1) / 2.e-307 = 2/2.e-307
      = 1.e307
    f = (b - a * r) / den = (1 - 1 * 1) / 2.e-307
      = 0

```

We see that, while the naive division generated an underflow, Smith's method produces the correct result.

Now that we have designed a more robust algorithm, we are interested in testing Smith's method on a more difficult case.

4.4 One more step

In this section, we show the limitations of Smith's method and present an example where Smith's method does not perform as expected.

The following example is inspired by an example by Stewart's in [42]. While Stewart gives an example based on a machine with an exponent range ± 99 , we consider an example which is based on Scilab's doubles. Consider the complex division

$$\frac{10^{307} + i10^{-307}}{10^{204} + i10^{-204}} \approx 1.0000000000000000 \cdot 10^{103} - i1.0000000000000000 \cdot 10^{-305}, \quad (84)$$

which is accurate to the displayed digits. In fact, there are more than 100 zeros following the leading 1, so that the previous approximation is very accurate. The following Scilab session compares the naive implementation, Smith's method and Scilab's division operator. The session is performed with Scilab v5.2.0 under a 32 bits Windows using a Intel Xeon processor.

```
-->[e f] = naive ( 1.e307 , 1.e-307 , 1.e204 , 1.e-204 )
f =
    0.
e =
    Nan
-->[e f] = smith ( 1.e307 , 1.e-307 , 1.e204 , 1.e-204 )
f =
    0.
e =
    1.000+103i
-->(1.e307 + %i * 1.e-307)/(1.e204 + %i * 1.e-204)
ans =
    1.000+103i - 1.000-305i
```

In the previous case, the naive implementation does not produce any correct digit, as expected. Smith's method, produces a correct real part, but an inaccurate imaginary part. Once again, Scilab's division operator provides the correct answer.

We first check why the naive implementation is not accurate in this case. We have $a = 10^{307}$, $b = 10^{-307}$, $c = 10^{204}$ and $d = 10^{-204}$. Indeed, the naive implementation performs the following steps.

```
den = c * c + d * d = 1.e204 * 1.e204 + 1.e-204 * 1.e-204
    = Inf
e = (a * c + b * d) / den
    = (1.e307 * 1.e204 + 1.e-307 * 1.e-204) / Inf = Inf / Inf
    = Nan
f = (b * c - a * d) / den
    = (1.e-307 * 1.e204 - 1.e307 * 1.e-204) / Inf = -1.e103 / Inf
    = 0
```

We see that the denominator `den` is in overflow, which makes `e` to be computed as `Nan` and `f` to be computed as 0.

Second, we check that Smith's formula is not accurate in this case. Indeed, it performs the following steps.

```
if ( abs(d) = 1.e-204 <= abs(c) = 1.e204 ) > test true
r = d/c = 1.e-204 / 1.e204 = 0
den = c + d * r = 1.e204 + 0 * 1.e-204 = 1.e204
e = (a + b * r) / den = (1.e307 + 1.e-307 * 0) / 1e204
    = 1.e307 / 1.e204 = 1.e103
f = (b - a * r) / den = (1.e-307 - 1.e307 * 0) / 1e204
    = 1.e-307 / 1.e204 = 0
```

We see that the variable `r` is in underflow, so that it is represented by zero. This simplifies the denominator `den`, but this variable is still correctly computed, because it is dominated the term `c`. The real part `e` is still accurate, because, once again, the computation is dominated by the term `a`. The imaginary part `f` is wrong, because this term should be dominated by the term `a*r`. Since `r` is in underflow, it is represented by zero, which completely changes the result of the expression `b-a*r`,

Scilab v5.2.0 release	Windows 32 bits	1.000+103 - 1.000-305i
Scilab v5.2.0 release	Windows 64 bits	1.000+103
Scilab v5.2.0 debug	Windows 32 bits	1.000+103
Scilab v5.1.1 release	Windows 32 bits	1.000+103
Scilab v4.1.2 release	Windows 32 bits	1.000+103
Scilab v5.2.0 release	Linux 32 bits	1.000+103 - 1.000-305i
Scilab v5.1.1 release	Linux 32 bits	1.000+103 - 1.000-305i
Octave v3.0.3	Windows 32 bits	1.0000e+103
Matlab 2008	Windows 32 bits	1.0000e+103 -1.0000e-305i
Matlab 2008	Windows 64 bits	1.0000e+103
FreeMat v3.6	Windows 32 bits	1.0000e+103 -1.0000e-305i

Figure 1: Result of the complex division $(1.e307 + \%i * 1.e-307)/(1.e204 + \%i * 1.e-204)$ on various softwares and operating systems.

which is now equal to b . Therefore, the result is equal to $1.e-307 / 1.e204$, which underflows to zero.

Finally, we analyze why Scilab’s division operator performs accurately in this case. Indeed, the formula used by Scilab is based on Smith’s method and we proved that this method fails in this case, when we use double floating point numbers. Therefore, we experienced here an unexpected high accuracy.

We performed this particular complex division over several common computing systems such as various versions of Scilab, Octave, Matlab and FreeMat on various operating systems. The results are presented in figure 1. Notice that, on Matlab, Octave and FreeMat, the syntax is different and we used the expression $(1.e307 + i * 1.e-307)/(1.e204 + i * 1.e-204)$.

The reason of the discrepancies of the results is the following [33, 30]. The processor being used may offer an internal precision that is wider than the precision of the variables of a program. Indeed, processors of the IA32 architecture (Intel 386, 486, Pentium etc. and compatibles) feature a floating-point unit often known as "x87". This unit has 80-bit registers in "double extended" format with a 64-bit mantissa and a 15-bit exponent. The most usual way of generating code for the IA32 is to hold temporaries - and, in optimized code, program variables - in the x87 registers. Hence, the final result of the computations depend on how the compiler allocates registers. Since the double extended format of the x87 unit uses 15 bits for the exponent, it can store floating point numbers associated with binary exponents from $2^{-16382} \approx 10^{-4932}$ up to $2^{16383} \approx 10^{4931}$, which is much larger than the exponents from the 64-bits double precision floating point numbers (ranging from $2^{-1022} \approx 10^{-308}$ up to $2^{1023} \approx 10^{307}$). Therefore, the computations performed with the x87 unit are less likely to generate underflows and overflows. On the other hand, SSE2 extensions introduced one 128-bit packed floating-point data type. This 128-bit data type consists of two IEEE 64-bit double-precision floating-point values packed into a double quadword.

Depending on the compilers options used to generate the binary, the result may use either the x87 unit (with 80-bits registers) or the SSE unit. Under Windows

32 bits, Scilab v5.2.0 is compiled with the `"/arch:IA32"` option [9], which allows Scilab to run on older Pentium computers that does not support SSE2. In this situation, Scilab may use the x87 unit. Under Windows 64 bits, Scilab uses the SSE2 unit so that the result is based on double precision floating point numbers only. Under Linux, Scilab is compiled with gcc [13], where the behavior is driven by the `-mfpmath` option. The default value of this option for i386 machines is to use the 387 floating point co-processor while, for x86_64 machines, the default is to use the SSE instruction set.

4.5 References

The 1962 paper by R. Smith [40] describes the algorithm which is used in Scilab.

Goldberg introduces in [18] many of the subjects presented in this document, including the problem of the complex division.

An analysis of Hough, cited by Coonen [8] and Stewart [42] shows that when the algorithm works, it returns a computed value \bar{z} satisfying

$$|\bar{z} - z| \leq \epsilon |z|, \quad (85)$$

where z is the exact complex division result and ϵ is of the same order of magnitude as the rounding unit for the arithmetic in question.

The limits of Smith's method have been analyzed by Stewart's in [42]. The paper separates the relative error of the complex numbers and the relative error made on real and imaginary parts. Stewart's algorithm is based on a theorem which states that if $x_1 \dots x_n$ are n floating point representable numbers, and if their product is also a representable floating point number, then the product $\min_{i=1,n}(x_i) \cdot \max_{i=1,n}(x_i)$ is also representable. The algorithm uses that theorem to perform a correct computation.

Stewart's algorithm is superseded by the one by Li et Al [27], but also by Kahan's [23], which, from [36], is the one implemented in the C99 standard.

Knuth presents in [26] the Smith's method in section 4.2.1, as exercise 16. Knuth gives also references [44] and [15]. The 1967 paper by Friedland [15] describes two algorithms to compute the absolute value of a complex number $|x + iy| = \sqrt{x^2 + y^2}$ and the square root of a complex number $\sqrt{x + iy}$.

Issues related to the use of extended double precision floating point numbers are analyzed by Muller et al. in [33]. In the section 3 of part I, "Floating point formats an Environment", the authors analyze the "double rounding" problem which occurs when an internal precision is wider than the precision of the variables of a program. The typical example is the double-extended format available on Intel platforms. Muller et al. show different examples, where the result depends on the compiler options and the platform, including an example extracted from a paper by Monniaux [30].

Corden and Kreitzer analyse in [9] the effect of the Intel compiler floating point options on the numerical results. The paper focuses on the reproducibility issues which are associated with floating point computations. The options which allow to be compliant with the IEEE standards for C++ and Fortran are presented. The

effects of optimization options is considered with respect to speed and the safety of the transformations that may be done on the source code.

The "Intel 64 and IA-32 Architectures Software Developer's Manual. Volume 1: Basic Architecture" [10] is part of a set of documents that describes the architecture and programming environment of Intel 64 and IA-32 architecture processors. The chapter 8, "Programming with the x87 environment" presents the registers and the instruction set for this unit. The section 8.1.2, "x87 FPU Data Registers" focuses on the floating point registers, which are based on 80-bits and implements the double extended-precision floating-point format. The chapter 10, "Programming with Streaming SIMD Extensions (SSE)" introduces the extensions which were introduced into the IA-32 architecture in the Pentium III processor family. The chapter 11 introduces the SSE2 extensions.

In [30], David Monniaux presents issues related to the analysis of floating point programs. He emphasizes the difficulty of defining the semantics of common implementation of floating point numbers, depending on choices made by the compiler. He gives concrete examples of problems that can appear and solutions.

5 Conclusion

We have presented several cases where the mathematically perfect algorithm (i.e. without obvious bugs) does not produce accurate results with the computer in particular situations. Many Scilab algorithms take floating point values as inputs, and return floating point values as output. We have presented situations where the intermediate calculations involve terms which are not representable as floating point values. We have also presented examples where cancellation occurs so that the rounding errors dominate the result. We have analyzed specific algorithms which can be used to cure some of the problems.

Most algorithms provided by Scilab are designed specifically to take into account for floating point numbers issues. The result is a collection of robust algorithms which, most of the time, exceed the user's needs.

Still, it may happen that the algorithm used by Scilab is not accurate enough, so that floating point issues may occur in particular cases. We cannot pretend that Scilab always use the best algorithm. In fact, we have given in this document practical (but extreme) examples where the algorithm used by Scilab is not accurate. In this situation, an interesting point is that Scilab is open-source, so that anyone who want can inspect the source code, analyze the algorithm and point out the problems of this algorithm.

That article does not aim at discouraging from using floating point numbers or implementing our own algorithms. Instead, the goal of this document is to give examples where some specific work is to do when we translate the mathematical material into a computational algorithm based on floating point numbers. Indeed, accuracy can be obtained with floating point numbers, provided that we are less *naïve*, use the appropriate theory and algorithms, and perform the computations with tested softwares.

6 Acknowledgments

I would like to thank Bruno Pinçon who made many highly valuable numerical comments on the section devoted to numerical derivatives. I would like to thank Claude Gomez for his support during the writing of this document. I would like to thank Bernard Hugueney and Allan Cornet who helped me in the analysis of the complex division portability issues.

7 Appendix

In this section, we analyze the examples given in the introduction of this article. In the first section, we analyze how the real number 0.1 is represented by a double precision floating point number, which leads to a rounding error. In the second section, we analyze how the computation of $\sin(\pi)$ is performed. In the final section, we make an experiment which shows that $\sin(2^{10i}\pi)$ can be arbitrary far from zero when we compute it with double precision floating point numbers.

7.1 Why 0.1 is rounded

In this section, we present a brief explanation for the following Scilab session.

```
-->format(25)
-->x1=0.1
x1   =
    0.10000000000000000055511
-->x2 = 1.0-0.9
x2   =
    0.09999999999999999777955
-->x1==x2
ans  =
    F
```

We see that the real decimal number 0.1 is displayed as 0.100000000000000005. In fact, only the 17 first digits after the decimal point are significant : the last digits are a consequence of the approximate conversion from the internal binary double number to the decimal number.

In order to understand what happens, we must decompose the floating point number into its binary components. The IEEE double precision floating point numbers used by Scilab are associated with a radix (or basis) $\beta = 2$, a precision $p = 53$, a minimum exponent $e_{min} = -1023$ and a maximum exponent $e_{max} = 1024$. Any floating point number x is represented as

$$fl(x) = M \cdot \beta^{e-p+1}, \quad (86)$$

where

- e is an integer called the exponent,
- M is an integer called the integral significant.

The exponent satisfies $e_{min} \leq e \leq e_{max}$ while the integral significant satisfies $|M| \leq \beta^p - 1$.

Let us compute the exponent and the integral significant of the number $x = 0.1$. The exponent is easily computed by the formula

$$e = \lfloor \log_2(|x|) \rfloor, \quad (87)$$

where the \log_2 function is the base-2 logarithm function. In the case where an underflow or an overflow occurs, the value of e is restricted into the minimum and maximum exponents range. The following session shows that the binary exponent associated with the floating point number 0.1 is -4.

```
-->format(25)
-->x = 0.1
x =
    0.1000000000000000055511
-->e = floor(log2(x))
e =
    - 4.
```

We can now compute the integral significant associated with this number, as in the following session.

```
-->M = x/2^(e-p+1)
M =
    7205759403792794.
```

Therefore, we deduce that the integral significant is equal to the decimal integer $M = 7205759403792794$. This number can be represented in binary form as the 53 binary digit number

$$M = 11001100110011001100110011001100110011001100110011010. \quad (88)$$

We see that a pattern, made of pairs of 11 and 00 appears. Indeed, the real value 0.1 is approximated by the following infinite binary decomposition:

$$0.1 = \left(\frac{1}{2^0} + \frac{1}{2^1} + \frac{0}{2^2} + \frac{0}{2^3} + \frac{1}{2^4} + \frac{1}{2^5} + \dots \right) \cdot 2^{-4}. \quad (89)$$

We see that the decimal representation of $x = 0.1$ is made of a finite number of digits while the binary floating point representation is made of an infinite sequence of digits. But the double precision floating point format must represent this number with 53 bits only.

Notice that, the first digit is not stored in the binary double format, since it is assumed that the number is *normalized* (that is, the first digit is assumed to be one). Hence, the leading binary digit is *implicit*. This is why there is only 52 bits in the mantissa, while we use 53 bits for the precision p . For the sake of simplicity, we do not consider denormalized numbers in this discussion.

In order to analyze how the rounding works, we look more carefully to the integer M , as in the following experiments.

```
-->7205759403792793 * 2^(-4-53+1)
ans =
```

```

0.0999999999999999916733

-->7205759403792794 * 2^(-4-53+1)
ans =
0.10000000000000000055511

```

We see that the real number 0.1 is between two consecutive floating point numbers:

$$7205759403792793 \cdot 2^{-4-53+1} < 0.1 < 7205759403792794 \cdot 2^{-4-53+1}. \quad (90)$$

There are four rounding modes in the IEEE floating point standard. The default rounding mode is *round to nearest*, which rounds to the nearest floating point number. In case of a tie, the rounding is performed to the only one of these two consecutive floating point numbers whose integral significant is even. In our case, the distance from the real x to the two floating point numbers is

$$|0.1 - 7205759403792793 \cdot 2^{-4-53+1}| = 8.33 \cdots 10^{-18}, \quad (91)$$

$$|0.1 - 7205759403792794 \cdot 2^{-4-53+1}| = 5.55 \cdots 10^{-18}. \quad (92)$$

(The previous computation is performed with a symbolic computation system, not with Scilab). Therefore, the nearest is the second integral significant. This is why the integral significant M associated with $x = 0.1$ is equal to 7205759403792794, which leads to $fl(0.1) \approx 0.100000000000000005$.

On the other side, $x = 0.9$ is also not representable as an exact binary floating point number (but 1.0 is exactly represented). The floating point binary representation of $x = 0.9$ is associated with the exponent $e = -1$ and an integral significant between 8106479329266892 and 8106479329266893. The integral significant which is nearest to $x = 0.9$ is 8106479329266893, which is associated with the approximated decimal number $fl(0.9) \approx 0.900000000000000002$.

Then, when we perform the subtraction "1.0-0.9", the decimal representation of the result is $fl(1.0) - fl(0.9) \approx 0.09999999999999997$, which is different from $fl(0.1) \approx 0.100000000000000005$.

7.2 Why $\sin(\pi)$ is rounded

In this section, we present a brief explanation of the following Scilab 5.1 session, where the function sinus is applied to the number π .

```

-->format(10)
ans =
0.
-->sin(%pi)
ans =
1.225D-16

```

This article is too short to make a complete presentation of the computation of elementary functions. The interested reader may consider the direct analysis of the Fdlibm library as very instructive [43]. Muller presents in "Elementary Functions" [32] a complete discussion on this subject.

In Scilab, the sin function is connected to a fortran source code (located in the *sci_f_sin.f* file), where we find the following algorithm:


```

do i = 0 , mn - 1
    y(i) = sin(x(i))
enddo

```

The `mn` variable contains the number of elements in the matrix, which is stored as the raw array `x`. This implies that no additional algorithm is performed directly by Scilab and the `sin` function is computed by the mathematical library provided by the compiler, i.e. by `gcc` under Linux and by Intel's Visual Fortran under Windows.

Let us now analyze the algorithm which is performed by the mathematical library providing the `sin` function. In general, the main structure of these algorithms is the following:

- scale the input x so that it lies in a restricted interval,
- use a polynomial approximation of the local behavior of `sin` in the neighborhood of 0.

In the `Fdlibm` library for example, the scaling interval is $[-\pi/4, \pi/4]$. The polynomial approximation of the `sin` function has the general form

$$\sin(x) \approx x + a_3x^3 + \dots + a_{2n+1}x^{2n+1} \quad (93)$$

$$\approx x + x^3p(x^2) \quad (94)$$

In the `Fdlibm` library, 6 terms are used.

For the `atan` function, which is used to compute an approximated value of π , the process is the same. This leads to a rounding error in the representation of π which is computed by Scilab as `4 * atan(1.0)`. All these operations are guaranteed with some precision, when applied to a number in the scaled interval. For inputs outside the scaling interval, the accuracy depends on the algorithm used for the scaling.

All in all, the sources of errors in the floating point computation of `sin(π)` are the following

- the error of representation of π ,
- the error in the scaling,
- the error in the polynomial representation of the function `sin`.

Since the value of `sin(π)` is close to the machine epsilon corresponding to IEEE double precision floating point numbers (i.e. close to $\epsilon \approx 2.220 \cdot 10^{-16}$), we see that the result is the best possible.

7.3 One more step

In fact, it is possible to reduce the number of significant digits of the sine function to as low as 0 significant digits. We mathematically have $\sin(2^n\pi) = 0$, but this can be very inaccurate with floating point numbers. In the following Scilab session, we compute `sin($2^{10i}\pi$)` for $i = 1$ to 5.

```
-->sin(2.^(10*(1:5)).*%pi).'  
ans =  
- 0.00000000000001254038322  
- 0.0000000001284092832066  
- 0.0000001314911060035225  
- 0.0001346468921407542141  
- 0.1374419777062635961151
```

For $\sin(2^{50}\pi)$, the result is very far from being zero. This computation may sound *extreme*, but it must be noticed that it is inside the IEEE double precision range of values, since $2^{50} \approx 3.10^{15} \ll 10^{308}$. If accurate computations of the sin function are required for large values of x (which is rare in practice), the solution may be to use multiple precision floating point numbers, such as in the MPFR library [31, 14], based on the Gnu Multiple Precision library [17].

References

- [1] Loss of significance. http://en.wikipedia.org/wiki/Loss_of_significance.
- [2] Maxima. <http://maxima.sourceforge.net/>.
- [3] Octave. <http://www.gnu.org/software/octave>.
- [4] Quadratic equation. http://en.wikipedia.org/wiki/Quadratic_equation.
- [5] Quadratic equation. <http://mathworld.wolfram.com/QuadraticEquation.html>.
- [6] M. Abramowitz and I. A. Stegun. *Handbook of Mathematical Functions with Formulas, Graphs, and Mathematical Tables*. 1972.
- [7] The Scilab Consortium. Scilab. <http://www.scilab.org>.
- [8] J.T. Coonen. Underflow and the denormalized numbers. *Computer*, 14(3):75–87, March 1981.
- [9] Martyn J. Corden and David Kreitzer. Consistency of floating-point results using the intel compiler or why doesn't my application always give the same answer? Technical report, Intel Corporation, Software Solutions Group, 2009.
- [10] Intel Corporation. Intel 64 and ia-32 architectures software developer's manual. volume 1: Basic architecture. <http://www.intel.com/products/processor/manuals>, 2009.
- [11] J. Dumontet and J. Vignes. Détermination du pas optimal dans le calcul des dérivées sur ordinateur. *R.A.I.R.O Analyse numérique*, 11(1):13–25, 1977.
- [12] George E. Forsythe. How do you solve a quadratic equation ? 1966. <ftp://reports.stanford.edu/pub/cstr/reports/cs/tr/66/40/CS-TR-66-40.pdf>.

- [13] Free Software Foundation. The gnu compiler collection. Technical report, 2008.
- [14] Laurent Fousse, Guillaume Hanrot, Vincent Lefèvre, Patrick Pélissier, and Paul Zimmermann. MPFR: A multiple-precision binary floating-point library with correct rounding. *ACM Transactions on Mathematical Software*, 33(2):13:1–13:15, June 2007.
- [15] Paul Friedland. Algorithm 312: Absolute value and square root of a complex number. *Commun. ACM*, 10(10):665, 1967.
- [16] P. E. Gill, W. Murray, and M. H. Wright. *Practical optimization*. Academic Press, London, 1981.
- [17] GMP. Gnu multiple precision arithmetic library. <http://gmplib.org>.
- [18] David Goldberg. *What Every Computer Scientist Should Know About Floating-Point Arithmetic*. Association for Computing Machinery, Inc., March 1991. http://www.physics.ohio-state.edu/~dws/group/links/floating_point_math.pdf.
- [19] Nicholas J. Higham. *Accuracy and Stability of Numerical Algorithms*. Society for Industrial and Applied Mathematics, Philadelphia, PA, USA, second edition, 2002.
- [20] P. Dugac J. Dixmier. *Cours de Mathématiques du premier cycle, 1ère année*. Gauthier-Villars, 1969.
- [21] M. A. Jenkins. Algorithm 493: Zeros of a real polynomial [c2]. *ACM Trans. Math. Softw.*, 1(2):178–189, 1975.
- [22] M. A. Jenkins and J. F. Traub. A three-stage algorithm for real polynomials using quadratic iteration. *SIAM Journal on Numerical Analysis*, 7(4):545–566, 1970.
- [23] W. KAHAN. Branch cuts for complex elementary functions, or much ado about nothing’s sign bit. pages 165–211, 1987.
- [24] W. Kahan. On the cost of floating-point computation without extra-precise arithmetic. 2004. <http://www.cs.berkeley.edu/~wkahan/Qdrtcs.pdf>.
- [25] C. T. Kelley. *Solving nonlinear equations with Newton’s method*. SIAM, 2003.
- [26] D. E. Knuth. *The Art of Computer Programming, Volume 2, Seminumerical Algorithms*. Third Edition, Addison Wesley, Reading, MA, 1998.
- [27] Xiaoye S. Li, James W. Demmel, David H. Bailey, Greg Henry, Yozo Hida, Jimmy Iskandar, William Kahan, Suh Y. Kang, Anil Kapur, Michael C. Martin, Brandon J. Thompson, Teresa Tung, and Daniel J. Yoo. Design, implementation and testing of extended and mixed precision blas. *ACM Trans. Math. Softw.*, 28(2):152–205, 2002.

- [28] Maple. Maplesoft. <http://www.maplesoft.com>.
- [29] The MathWorks. Matlab. <http://www.mathworks.fr>.
- [30] David Monniaux. The pitfalls of verifying floating-point computations. *ACM Trans. Program. Lang. Syst.*, 30(3):1–41, 2008.
- [31] MPFR. The mpfr library. <http://www.mpfr.org>.
- [32] Jean-Michel Muller. *Elementary functions: algorithms and implementation*. Birkhauser Boston, Inc., Secaucus, NJ, USA, 1997.
- [33] Jean-Michel Muller, Nicolas Brisebarre, Florent de Dinechin, Claude-Pierre Jeannerod, Vincent Lefèvre, Guillaume Melquiond, Nathalie Revol, Damien Stehlé, and Serge Torres. *Handbook of Floating-Point Arithmetic*. Birkhäuser Boston, 2010. ACM G.1.0; G.1.2; G.4; B.2.0; B.2.4; F.2.1., ISBN 978-0-8176-4704-9.
- [34] Yves Nievergelt. How (not) to solve quadratic equations. *The College Mathematics Journal*, 34(2):90–104, 2003.
- [35] W. H. Press, Saul A. Teukolsky, William T. Vetterling, and Brian P. Flannery. *Numerical Recipes in C, Second Edition*. 1992.
- [36] Douglas M. Priest. Efficient scaling for complex division. *ACM Trans. Math. Softw.*, 30(4):389–401, 2004.
- [37] Wolfram Research. Mathematica. <http://www.wolfram.com/products/mathematica>.
- [38] Wolfram Research. Wolfram alpha. <http://www.wolframalpha.com>.
- [39] K.E. Schmidt. Numerical derivatives. <http://fermi.la.asu.edu/PHY531/intro/node1.html>.
- [40] Robert L. Smith. Algorithm 116: Complex division. *Commun. ACM*, 5(8):435, 1962.
- [41] R. S. Stepleman and N. D. Winarsky. Adaptive numerical differentiation. *Mathematics of Computation*, 33(148):1257–1264, 1979.
- [42] G. W. Stewart. A note on complex division. *ACM Trans. Math. Softw.*, 11(3):238–241, 1985.
- [43] Inc. Sun Microsystems. A freely distributable c math library. 1993. <http://www.netlib.org/fdlibm>.
- [44] P. Wynn. An arsenal of ALGOL procedures for complex arithmetic. *BIT Numerical Mathematics*, 2(4):232–255, December 1962.

Index

derivative, [15](#)

numdiff, [19](#)

roots, [6](#)

absolute error, [3](#)

cancellation, [6](#)

Corden, Martyn, [28](#)

floating point numbers, [3](#)

Forsythe, George, [13](#)

Goldberg, David, [13](#), [28](#)

IEEE 754, [3](#)

Jenkins, M. A., [8](#)

Kahan, William, [13](#), [28](#)

Knuth, Donald E., [28](#)

Kreitzer, David, [28](#)

massive cancellation, [6](#)

Monniaux, David, [29](#)

Muller, Jean-Michel, [28](#)

overflow, [7](#)

relative error, [3](#)

Smith, Robert, [24](#), [28](#)

Stewart, G. W., [28](#)

Traub, J. F., [8](#)